Microsoft

Microsoft Defender Advanced Threat Protection

Attack simulation

Scenario 2: PowerShell script in fileless attack

February 2019

Copyright

Our detection philosophy

It's simple.

We make sure that known advanced persistent threat (APT) indicators or techniques are visible in our telemetry, that we recognize them, and that we are able to raise the relevant alerts.

When we raise an alert near real-time, we provide the relevant context, including actor attribution, their victimology, geo-affinity, and main tactics. This is realized through a rich, dynamic library of known attack indicators, including known threat components previously observed on real machines, script and web page snippets from compromised or malicious websites, as well as IPs, URLs, and domains representing the attacker's infrastructure. We constantly update this library with new threat intelligence generated mainly by Microsoft's own APT hunting and research teams, but enriched by collaboration with partners and shared feeds.

Because threats are constantly being crafted and modified, we monitor a large set of anomalous and suspicious behaviors to find new and unknown actor activity. These anomalous and suspicious activities raise alerts for the Security Operations Center (SOC) analyst to validate and address. With the help of information about proximate events on the same machine and other relevant machines, SOC analysts can validate actual breach activity, determine risk, establish the scope of the breach, define containment activities, and then contain, mitigate and fully respond to the attack.

Introduction: PowerShell script in fileless attack

In this scenario, we move up a notch to more sophisticated attack that leverage advanced techniques to stay under the detection radar. This category of attacks usually doesn't include files dropped on the victim's machine—they occur solely in memory. They "live off the land" by using only existing system and administrative tools and injecting their code into system processes to hide their execution and persist on the box.

We will simulate such an attack and explore how exploit protection capabilities in Windows 10 can help prevent attackers from being able to carry out some of their activities.

In this simulation, our example scenario starts with a PowerShell script. A user may be tricked into executing such a script, or the script may be executed remotely from another machine in the organization that was previously infected, with the attacker attempting to move laterally in the network. Detection of such scripts is difficult because administrators also often run scripts remotely to carry out various administrative activities.

During the simulation, the attacker goes on to inject some shellcode into a seemingly innocent process, in this case *notepad.exe*. We chose this process for the simulation, but attackers will more likely target a long-running system process like *svchost.exe*. The shellcode then goes on to contact the attacker's command-and-control (C&C) server to receive instructions on how to proceed.

The test machine require for this simulation should:

- Be onboarded to Microsoft Defender ATP
- Run Windows 10 Fall Creators Update (version 1709)
- Have PowerShell turned on
- Have Windows Defender Antivirus turned on

For onboarding instructions, read to the product guide. We recommend running the local onboarding script to onboard the test machine.

Run the simulation

To run this attack scenario, follow these steps:

1. On the designated test machine, log in to the Microsoft Defender ATP portal and go to **Help (?) > Simulations & tutorials**.

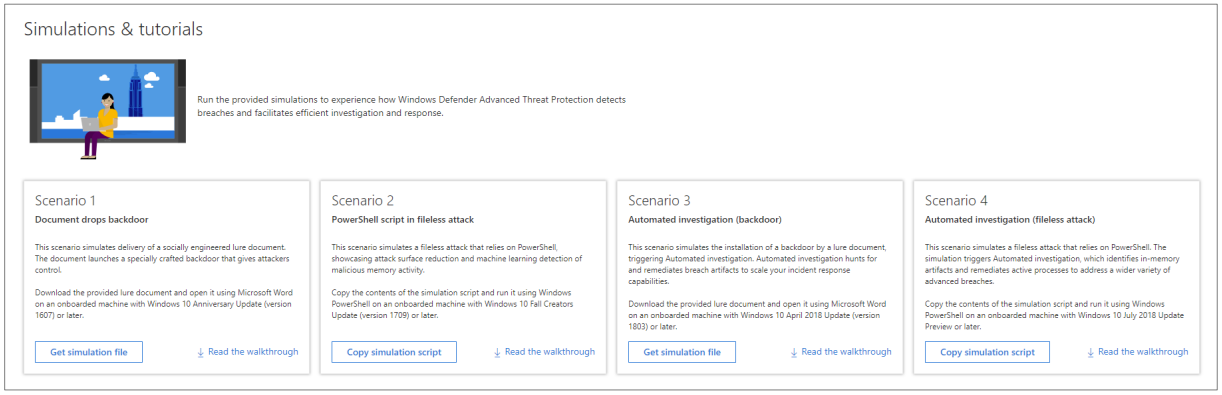

Figure 1. Simulation scenarios in the portal

2. Click the **Copy simulation script** button under **Scenario 2: PowerShell script in fileless attack** to copy the PowerShell script.
3. Open a Windows PowerShell window with administrative privileges on the test machine.
4. At the prompt, paste and run the provided script.

   A few seconds later, *notepad.exe* is started and the simulated attack code is injected into it. The simulated attack code attempts communication to an external IP address simulating the C&C server.

Simulate the attack with exploit protection

With exploit protection introduced with Windows 10 Fall Creators Update (version 1709), policies can be applied to restrict how code runs on machines, mitigating many exploit-based attacks. Exploit protection detections are surfaced as alerts by Microsoft Defender ATP to provide SOC personnel with visibility into these events.

In this section, we will configure *exploit protection* so that it disallows dynamic code execution in our process of interest, *notepad.exe*, and then run the simulated attack again.

To simulate the attack with exploit protection:

1. Open a Windows PowerShell window with administrative privileges.
2. At the prompt, run the following commands to configure exploit protection:

```
$path = "HKLM:\Software\Microsoft\Windows NT\CurrentVersion\Image File Execution Options\notepad.exe";
$value = ([byte[]](0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x11,0x11,0x01,0x01,0x00,0x00));
New-Item -Path $path -Force;
New-ItemProperty -Path $path -Name "MitigationOptions" -Value $value -PropertyType Binary -Force
```

   Note: The exploit protection configuration provided here is solely to illustrate pertinent functionality. This configuration should not be applied to other machines in production without proper analysis of its impact.

3. Now run the provided attack PowerShell script from the **Simulations & tutorials** page again.

   The script starts *notepad.exe* again, injects its malicious shellcode into it, and attempts to execute it as before. This time, however, it is stopped by exploit protection, which causes *notepad.exe* to be terminated.

4. [Optional] To restore exploit protection settings on the test machine, run the following command in the PowerShell window:

```
Remove-ItemProperty -Path $path -Name "MitigationOptions" -Force
```

Congrats – you're done running the attack!

The attack simulation ends here. A real attacker, if successful, would likely continue to scan for information, send collected reconnaissance information to a command-and-control (C&C) server, and use this information to move laterally and pursue other attractive targets.

Next, let's review and investigate the Microsoft Defender ATP alerts that surfaced the simulated attack.

Note: Alerts should start to appear 15-30 minutes after the simulated backdoor is launched.

Investigate the attack in the portal

Let's switch into our defender role and explore the attack from the SOC point of view in the Microsoft Defender ATP portal.

1. Open the Microsoft Defender ATP portal from any machine.
2. Log in with your Microsoft Defender ATP credentials. Default global administrator credentials are provided with your signup email.
3. After 15-30 minutes of the simulated attack, you should find several new alerts on the dashboard.

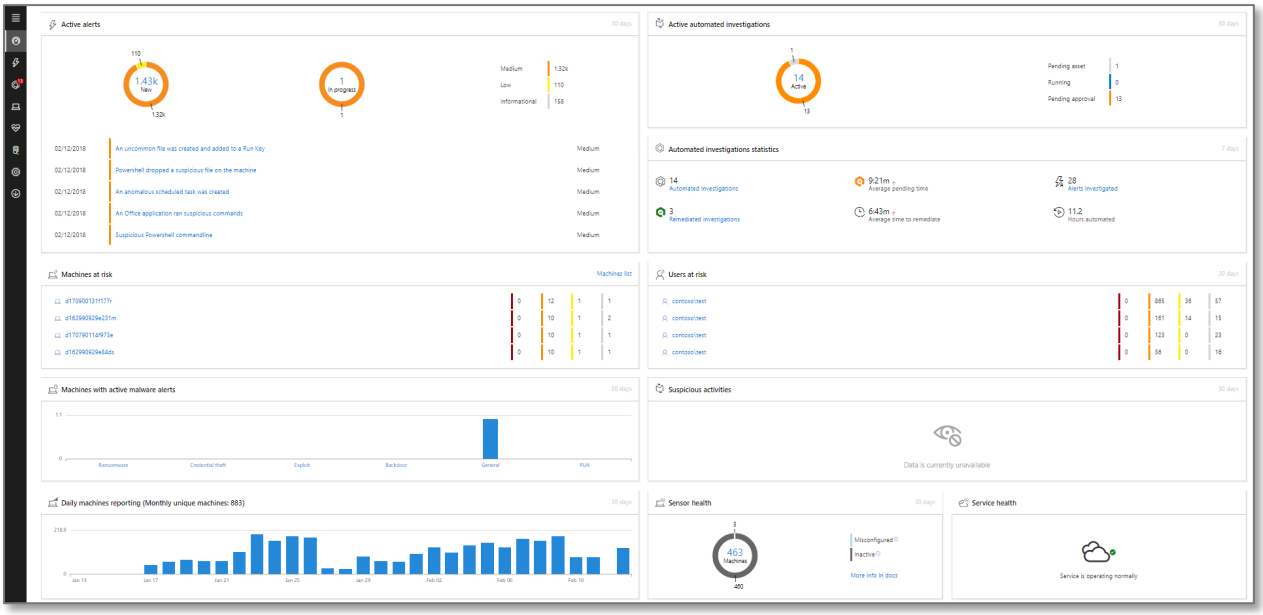

Figure 2. Dashboard view showing the alerts

Investigate the attack as a single incident

Microsoft Defender ATP applies correlation analytics and aggregates all related alerts and investigations into one "incident" entity. By doing so, Microsoft Defender ATP narrates a broader attack story, allowing the SOC analyst to understand and deal with complex threats across the org with the right visuals—through the enhanced incident graph—and data representations.

The alerts generated during this simulation are associated with the same threat, and as a result are automatically aggregated as a single incident.

To view the incident, go to the **Incidents** queue and select the relevant item as shown below. A side panel displays additional information about the incident, including all the related alerts.

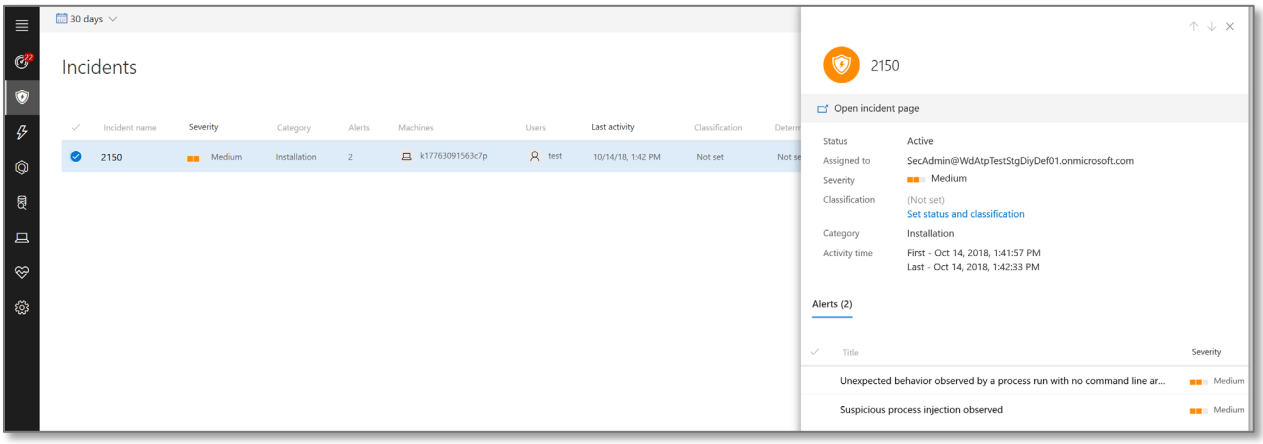

Figure 3. Incident aggregating alerts generated during the simulation

Select **Open incident page** to get more information about the incident.

In the incident page, you can check all the affected machines and the related alerts. For a broader view of the entities involved in the incident, select **Graph**.

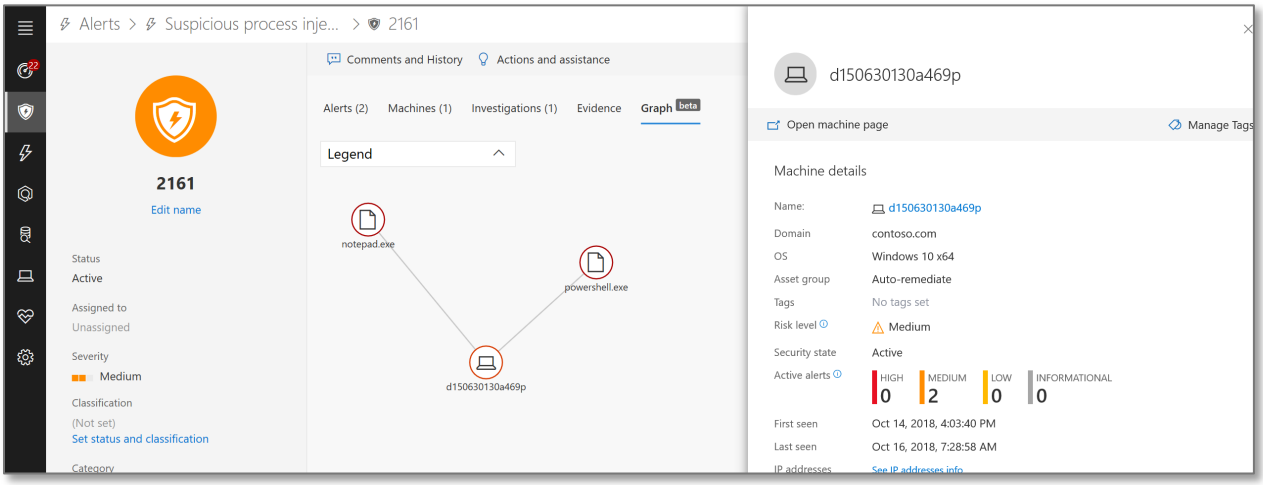

Figure 4. Graph of the incident

Reviewing the incident alert list unfolds the progression of the attack. From this view you can dive into the individual alerts

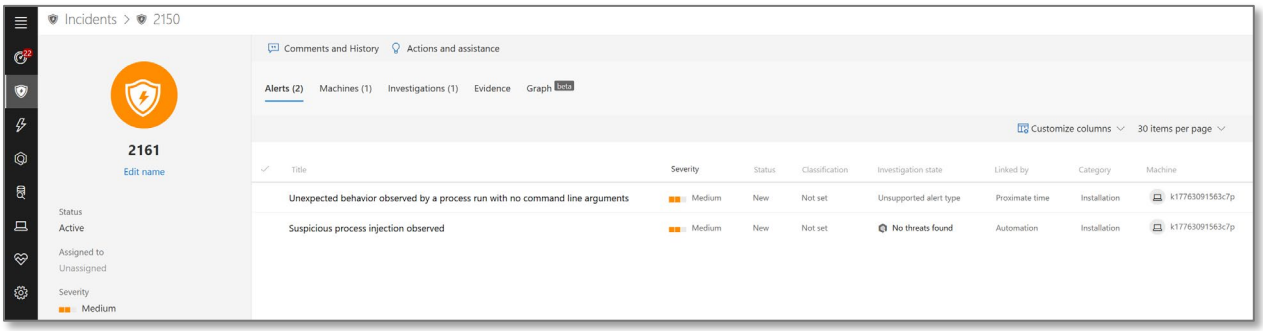

Figure 5. Incident related alerts

Review generated alerts

Let's look at some of the alerts generated during the simulated attack.

Note: We will walk through only a few of the alerts generated during the simulated attack. Depending on the version of Windows and the Windows Defender Antivirus protection updates running on your test machine, you might see more alerts and they might appear in a slightly different order.

Alert: Suspicious process injection observed

Advanced attackers will use more sophisticated and stealthy methods to persist in memory and hide from detection tools. One common technique is to operate from within a trusted system process rather than a malicious executable, making it hard for detection tools and security operations to spot the malicious code.

To allows SOC personnel to catch such advanced attacks, deep memory sensors in Microsoft Defender ATP provide our cloud service with unprecedented visibility into a variety of cross-process code injection techniques. As show below, Microsoft Defender ATP detected and alerted on the attempt to inject code to notepad.exe.

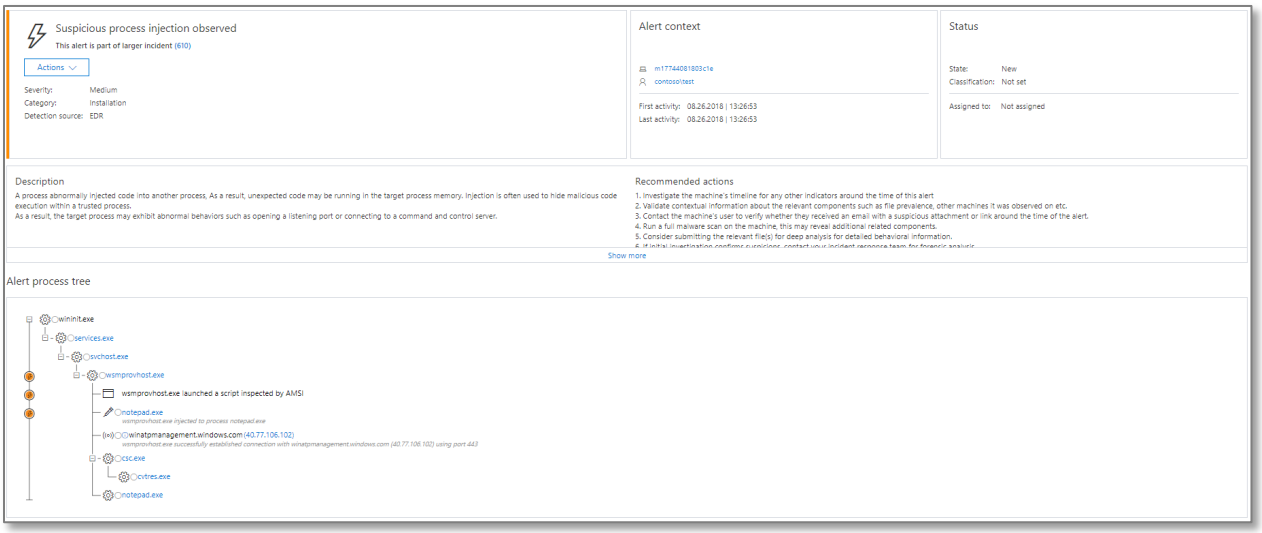

Figure 6. Alert for injection of potentially malicious code

Alert: Unexpected behavior observed by a process run with no command line arguments

Microsoft Defender ATP detections are often targeting the most invariant aspect of an attack technique. This ensures durability and raises the bar for attacker's to switch to newer tactics.

We employ large-scale learning algorithms to establish normal behavior of common processes within an organization and worldwide, and watch for when these processes exhibit anomalous behaviors. These anomalous behaviors often indicate that extraneous code was introduced and running in the otherwise trusted process.

In our case, the well-known process *notepad.exe* is exhibiting abnormal behavior, involving communication with an external location. Note that this outcome is independent of the specific method used to introduce and execute the malicious code.

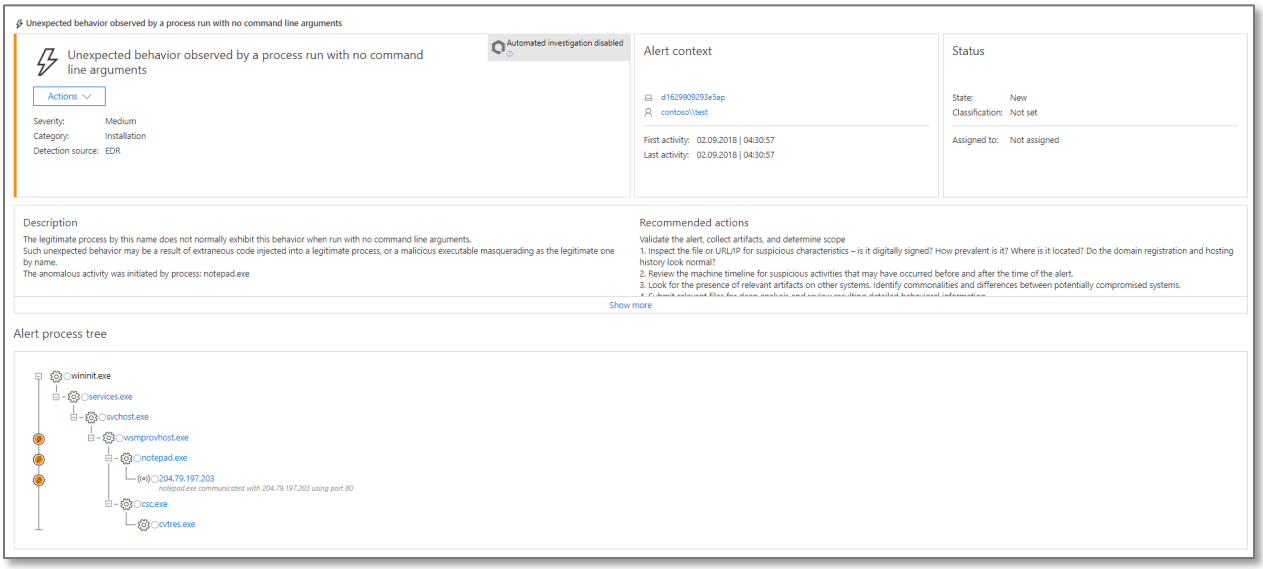

Figure 7. Alert for unexpected behavior by a process run with no command line arguments

✎ **Note**: Because this alert is based on machine-learning models that require some backend processing, it might take some time before it is actually generated on the portal.

Notice that the alert details include the external IP address—an indicator you can use as a pivot to expand investigation. Click the IP address in the **Alert Process Tree** to view the IP address details page.

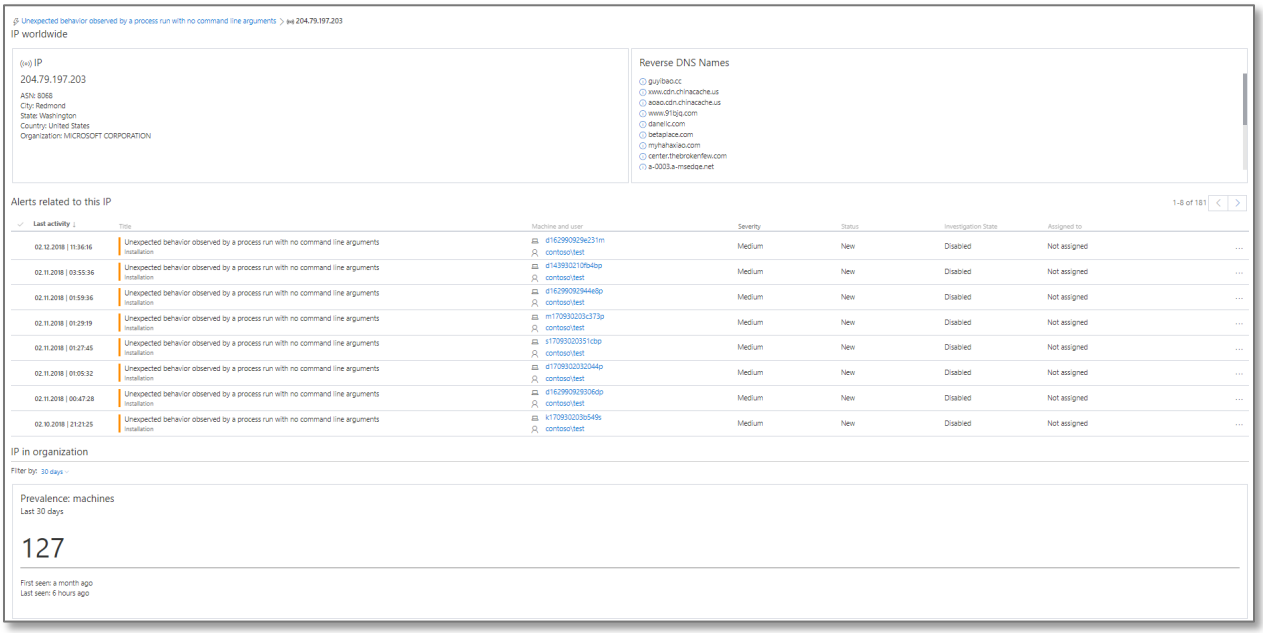

Figure 8. IP address details page

Alert: EAF violation blocked by exploit protection

By enabling exploit protection capabilities, we have hardened the victim machine against unexpected code execution. We have specifically enabled exploit protection rules that detect and prohibit unexpected code execution in *notepad.exe*. As a result, the simulated attempt to execute injected shellcode is detected and blocked, essentially stopping the attack's progression.

The exploit protection detection also results in a "EAF violation blocked by exploit protection" alert in the Microsoft Defender ATP portal. Note that the severity level of the alert is informational, because the attack was stopped. However, the security team is still notified that a possible exploitation attempt was made so they can take precautionary measures and be on the alert for a potentially persistent attacker.

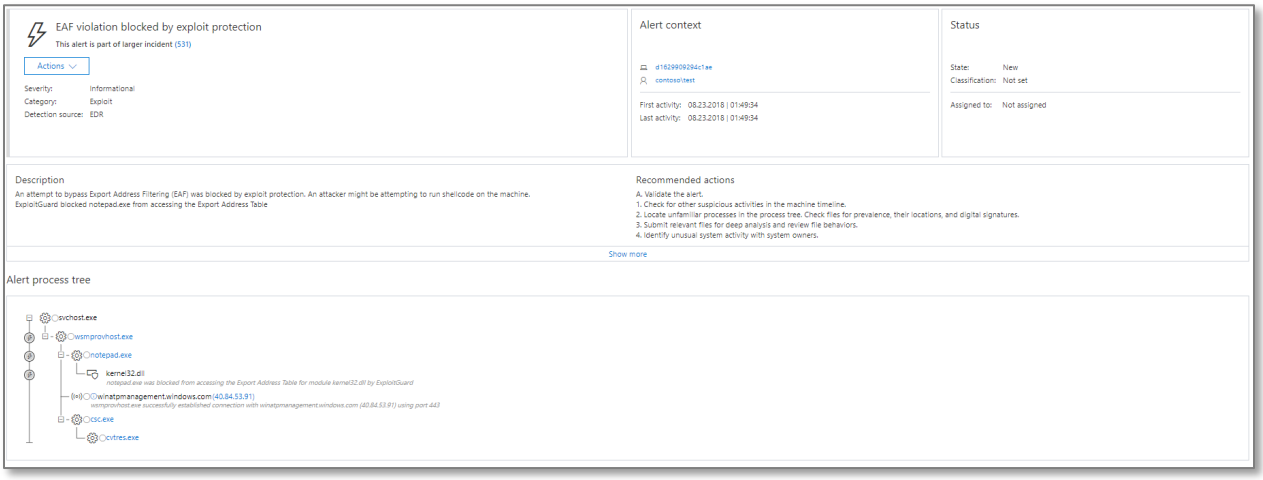

Figure 9. Alert for EAF violation detected by exploit protection

Review the machine timeline

Clicking on the machine name on one of the alert pages opens the machine details page. On this page, the alert itself and related events on the machine are provided to ease investigation. You can scroll through the machine timeline and view all events and behaviors observed on the machine in chronological order, interspersed with the alerts raised. Note the different information levels available: *detections, behaviors, and verbose*.

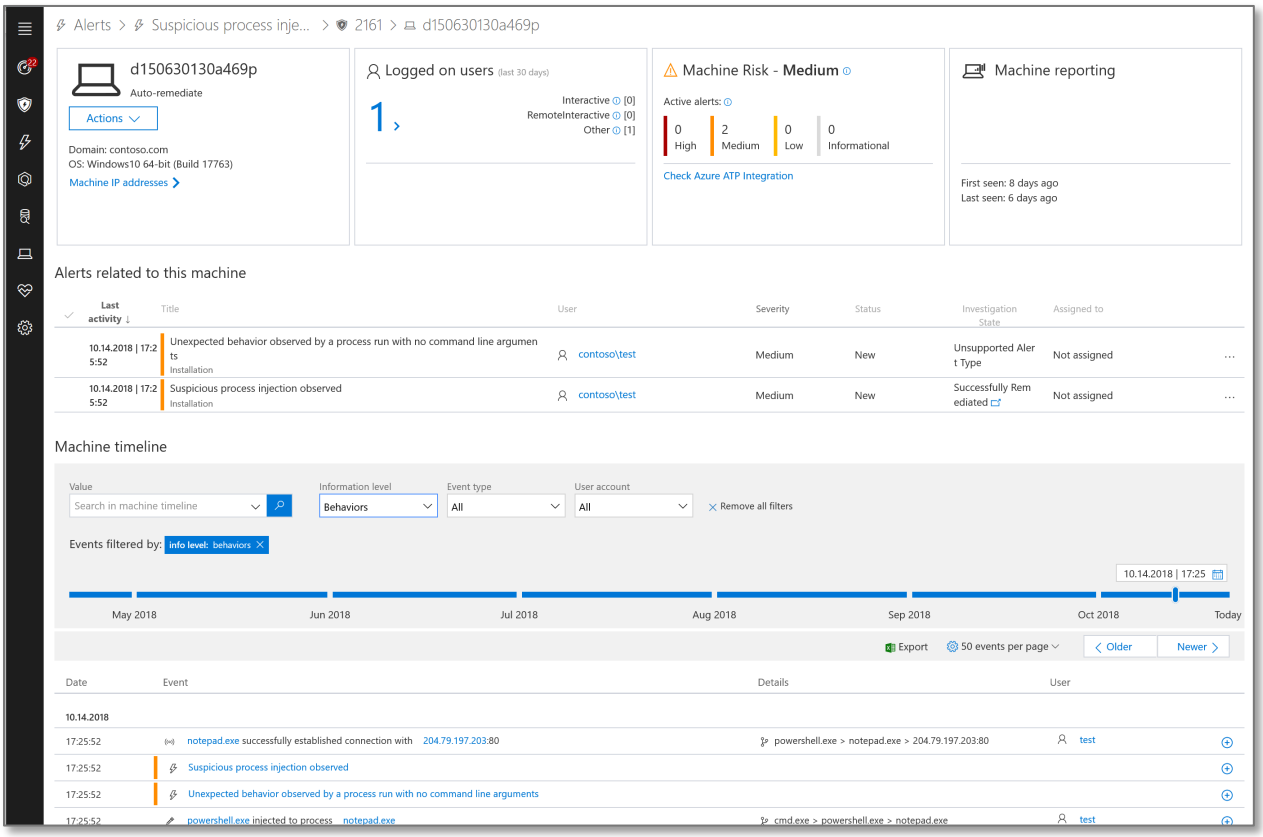

Figure 10. Machine timeline with behaviors

Expanding some of the more interesting behaviors provides useful details, such as process trees. For example, clicking on the item **powershell.exe injected to process notepad.exe** displays the full process tree for this behavior. Selecting the powershell.exe will show more information on this specific execution

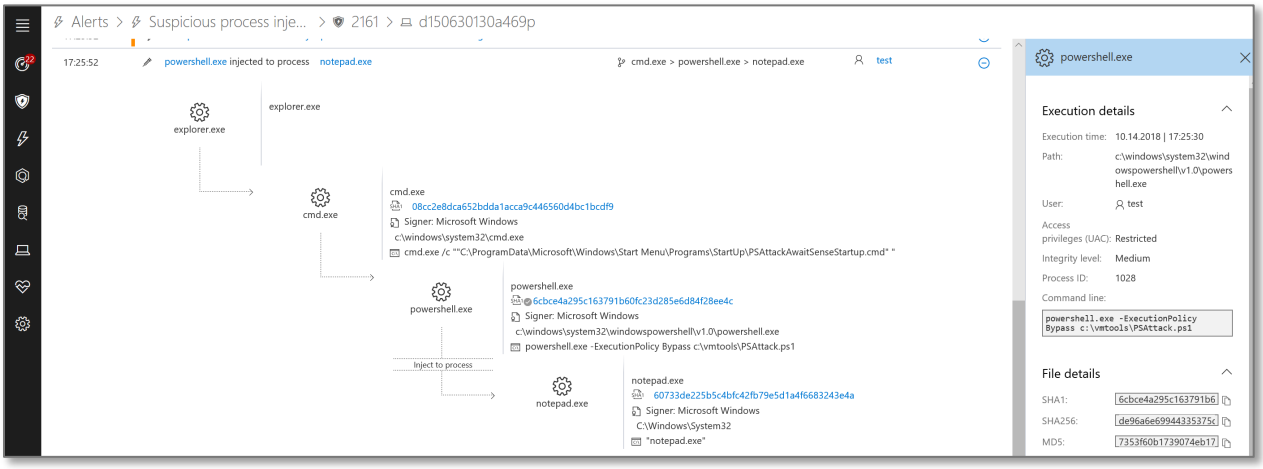

Figure 11. Process tree for selected PowerShell file creation behavior

Resolve the incident

Now that the investigation is completed and, in our case, confirmed to be a benign activity, it is time to close the incident.

On the incident page, select **Actions and assistance** to get management options that apply to the entire incident and all related alerts.

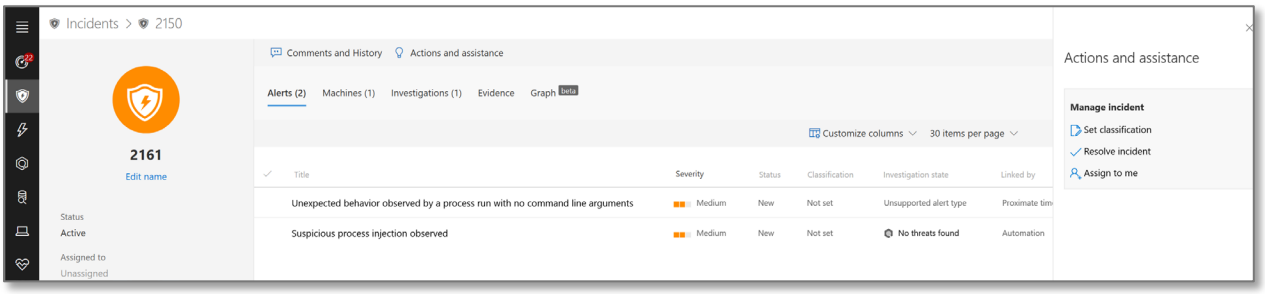

Figure 14. Resolving the incident and related alerts

Review the reports

Before concluding the investigation, it's a good idea to look at the reports dashboard. It provides high-level information about alerts and machine related information generated in your organization. The report includes trends and summary information on alerts and machines.

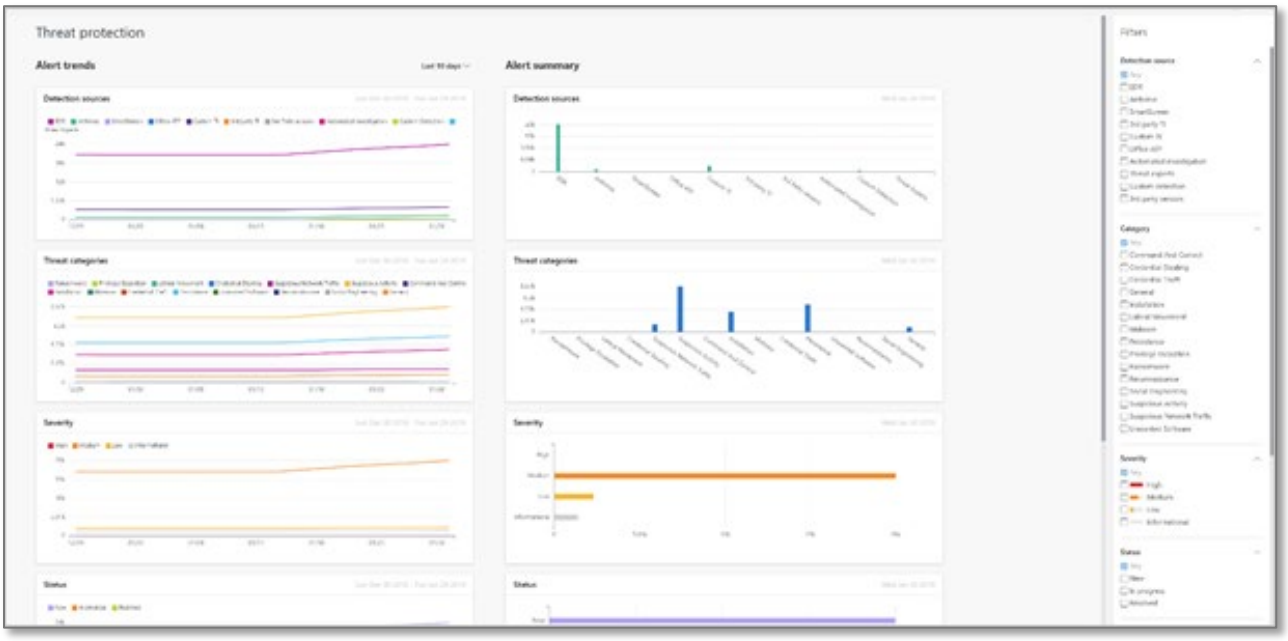

Figure 15. Threat protection report page

Knowing the trends and summaries related to alerts and machines in your organization can help identify where focused improvements can be made. For example, if you see a sudden spike in a specific kind of alert, you can drill down and start investigating directly from the relevant card to pivot into the alert or machine queue with the relevant filters applied and determine what action to take to address an issue.

Conclusion

We've simulated an advanced memory-only attack, and walked through how Microsoft Defender ATP detects and alerts on stealthy malicious activity with the help of deep OS sensors. We also experienced how exploit protection capabilities can stop advanced attacks and provide alert information in the portal. We've seen how alerts are delivered along with other contextual information, enabling SOC personnel to investigate and take necessary action.

We hope you enjoyed this simulation and are now encouraged to explore other features and capabilities. For more information, read the product guide at docs.microsoft.com.

Click the feedback icon on the Microsoft Defender ATP portal to let us know how you feel about this simulation or any other aspects of the product. We would love to hear your ideas about additional simulations and tutorials. Thank you!